

Programming  
Natural Language Processing  
Computer Vision  
News  
SQL


[JOIN NEWSLETTER](#)

Datasets  
Courses  
Online  
Certificates  
Webinars



[Trifacta](#) | [Transform Your Data in 6 Steps](#) | [Get ebook](#)

# Who is your Golden Goose?: Cohort Analysis

Step-by-step tutorial on how to perform customer segmentation using RFM analysis and K-Means clustering in Python.

[Share 16](#)

[Tweet](#)

[Share](#)

[Share](#)

9

Pages: 1 2

## K-Means Clustering

**K-Means clustering** is one type of unsupervised learning algorithms, which makes groups based on the distance between the points. How? There are two concepts of distance in K-Means clustering. **Within Cluster Sums of Squares**(WSS) and **Between Cluster Sums of Squares** (BSS).

- **Within Cluster Sums of Squares :** 
$$WSS = \sum_{i=1}^{N_c} \sum_{x \in C_i} d(x, \bar{x}_{C_i})^2$$
- **Between Cluster Sums of Squares:** 
$$BSS = \sum_{i=1}^{N_c} |C_i| \cdot d(\bar{x}_{C_i}, \bar{x})^2$$

$C_i$  = Cluster,  $N_c$  = # clusters,  $\bar{x}_{C_i}$  = Cluster centroid,  $\bar{x}$  = Sample Mean

WSS means the sum of distances between the points and the corresponding centroids for each cluster and BSS means the sum of distances between the centroids and the total sample mean multiplied by the number of points within each cluster. So you can consider WSS as the measure of compactness and BSS as the measure of separation. For clustering to be successful, we need to get the lower WSS and the higher BSS.

By iterating and moving the cluster centroids, K-Means algorithm tries to get the optimized points of the centroid, which minimize the value of WSS and maximize the value of BSS. I won't go more in-depth with the basic concept, but you can find a further explanation from video.

Search KDnuggets...



[Anaconda Nucleus](#)  
[Exclusive tutorials](#)  
[Interactive notebooks](#)



## Latest News

What's Your Passion? Make It a Reality During Our Cha...

5 Ways to Apply AI to Small Data Sets

Decision Tree Algorithm, Explained

Building a Visual Search Engine – Part 1: Data Ex...

KDnuggets 22:n06, Feb 9: Data Science Programming Languages...

What certification are you adding to your toolkit in 2022?

## Top Posts Last Week

1 Is Data Science a Dying Career?

Programming  
Natural Language Processing  
Computer Vision  
News



JOIN NEWSLETTER





Photo from Wikipedia



- 2 Top Five SQL Window Functions You Should Know For Data Science Interviews
- 3 A Deep Look Into 13 Data Scientist Roles and Their Responsibilities
- 4 SQL Interview Questions for Experienced Professionals
- 5 Why Do Machine Learning Models Die In Silence?

Because K-means clustering uses the distance as the similarity factor, we need to scale the data. Suppose we have two different scales of features, say height and weight. Height is over 150cm and weight is below 100kg on average. So If we plot this data, the distance between the points will be highly dominated by height resulting in a biased analysis.

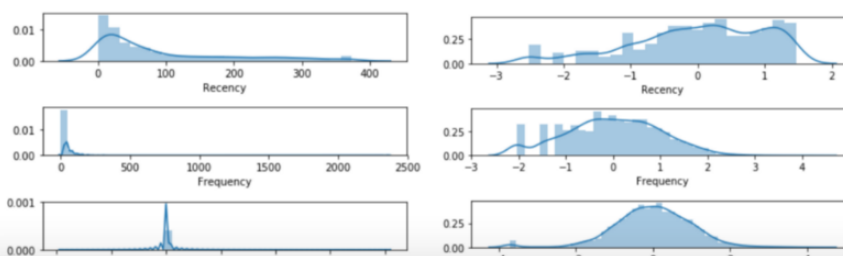
Therefore when it comes to K-means clustering, scaling and normalizing data is a critical step for preprocessing. If we check the distribution of RFM values, you can notice that they are right-skewed. It's not a good state to use without standardization. Let's transform the RFM values into log scaled first and then normalize them.

```
# define function for the values below 0
def neg_to_zero(x):
    if x <= 0:
        return 1
    else:
        return x
# apply the function to Recency and MonetaryValue column
rfm['Recency'] = [neg_to_zero(x) for x in rfm.Recency]
rfm['Monetary'] = [neg_to_zero(x) for x in rfm.Monetary]
# unskew the data
rfm_log = rfm[['Recency', 'Frequency', 'Monetary']].apply(np.log, axis = 1).ro
```

The values below or equal to zero go negative infinite when they are in log scale, I made a function to convert those values into 1 and applied it to Recency and Monetary column, using list comprehension like above. And then, a log transformation is applied for each RFM values. The next preprocessing step is scaling but it's simpler than the previous step.

Using **StandardScaler()**, we can get the standardized values like below.

```
# scale the data
scaler = StandardScaler()
rfm_scaled = scaler.fit_transform(rfm_log)
# transform into a dataframe
rfm_scaled = pd.DataFrame(rfm_scaled, index = rfm.index, columns = rfm_log.col
```



## More Recent Posts

What certification are you adding to your toolkit in 2022?

The Not-so-Sexy SQL Concepts to Make You Stand Out

Top Stories, Jan 31 – Feb 6: 7 Steps to Mastering Machin...

The Complete Collection of Data Science Cheat Sheets – P...

19 Data Science Project Ideas for Beginners

Build a Web Scraper with Python in 5 Minutes

Deploying a Streamlit WebApp to Heroku using DAGsHub

Data Science Programming Languages and When To Use Them

An Overview of Logistic Regression

8 Best Data Science Courses to Enroll in 2022 For Steep Career...

SHARES

Programming  
Natural Language Processing  
Computer Vision  
News



[JOIN NEWSLETTER](#)

easily and accurately. Now, we are done with preprocessing.

What is the next? The next step will be selecting the right number of clusters. We have to choose how many groups we're going to make. If there is prior knowledge, we can just give the number right ahead to the algorithm. But most of the case in unsupervised learning, there isn't. So we need to choose the optimized number, and the Elbow method is one of the solutions where we can get the hints.

```
# the Elbow method
wcss = {}
for k in range(1, 11):
    kmeans = KMeans(n_clusters= k, init= 'k-means++', max_iter= 300)
    kmeans.fit(rfm_scaled)
    wcss[k] = kmeans.inertia_
# plot the WCSS values
sns.pointplot(x = list(wcss.keys()), y = list(wcss.values()))
plt.xlabel('K Numbers')
plt.ylabel('WCSS')
plt.show()
```

Using for loop, I built the models for every number of clusters from 1 to 10. And then collect the WSS values for each model. Look at the plot below. As the number of clusters increases, the value of WSS decreases. There is no surprise cause the more clusters we make, the size of each cluster will decrease so the sum of the distances within each cluster will decrease. Then what is the optimal number?



The answer is at the 'Elbow' of this line. Somewhere WSS dramatically decrease but not too much K. My choice here is three. What do you say? Doesn't it really look like an elbow of the line?

Now we chose the number of clusters, we can build a model and make actual clusters like below. We can also check the distance between each point and the centroids or the labels

Build a Serverless News Data Pipeline using ML on AWS Cloud

Introducing dbt, the ETL and ELT Disrupter

How do we Better Solve Analytics Problems?

KDnuggets™ News 19:n40, Oct 23: How to Become a (Good) Data...

Hilarious Data Science Humor

The Uncommon Data Science Job Guide



A Two-Day Virtual ML Community Event for AI Developers and Data Scientists.

Get The Latest News!



Get KDnuggets, a leading newsletter on AI, Data Science and Machine Learning, straight to your inbox.

**SIGN UP**

By subscribing you accept KDnuggets Privacy Policy

SHARES

Programming  
Natural Language Processing  
Computer Vision  
News  
SQL


[JOIN NEWSLETTER](#)

▼ Datasets	Recency	Frequency	Monetary	R	F	M	RFM_Segment	RFM_Score	RFM_Level	K_Cluster	
CustomerID											
▼ Courses											
12347.0	Online	2	45	980.10	4	4	4	444	12.0	Gold	1
12348.0	Certificates	248	9	728.88	1	2	4	124	7.0	Bronze	0
12349.0	Webinars	18	25	486.87	3	3	3	333	9.0	Silver	0
▼ Resources											
12350.0		310	3	45.70	1	1	1	111	3.0	Green	2
12352.0	Meetings	36	32	1353.48	3	3	4	334	10.0	Silver	1
▼ Jobs											
Software											

Now we made two kinds of segmentation, RFM quantile groups and K-Means groups. Let's make visualization and compare the two methods.

### Snake plot and heatmap


I'm going to make two kinds of plot, a line plot and a heat map. We can easily compare the differences of RFM values with these two plots. Firstly, I'll make columns to assign the two clustering labels. And then reshape the data frame by melting the RFM values into one column.

```
# assign cluster column
rfm_scaled['K_Cluster'] = clus.labels_
rfm_scaled['RFM_Level'] = rfm.RFM_Level
rfm_scaled.reset_index(inplace = True)
# melt the dataframe
rfm_melted = pd.melt(frame= rfm_scaled, id_vars= ['CustomerID', 'RFM_Level', 'K_Cluster'])
rfm_melted.head()
```

	CustomerID	RFM_Level	K_Cluster	Metrics	Value
0	12347.0	Gold	1	Recency	-2.048810
1	12348.0	Bronze	0	Recency	1.193267
2	12349.0	Silver	0	Recency	-0.571042
3	12350.0	Green	2	Recency	1.343937
4	12352.0	Silver	1	Recency	-0.104236

This will make recency, frequency and monetary categories as observations, which allows us to plot the values in one plot. Put Metrics at x-axis and Value at y-axis and group the values by RFM\_Level. Repeat the same code which groups the values by K\_Cluster this time. The outcome would be like below.

```
# a snake plot with RFM
sns.lineplot(x = 'Metrics', y = 'Value', hue = 'RFM_Level', data = rfm_melted)
plt.title('Snake Plot of RFM')
plt.legend(loc = 'upper right')
# a snake plot with K-Means
sns.lineplot(x = 'Metrics', y = 'Value', hue = 'K_Cluster', data = rfm_melted)
plt.title('Snake Plot of RFM')
plt.legend(loc = 'upper right')
```




Let's try again with a heat map. Heat maps are a graphical representation of data where larger values were colored in darker scales and smaller values in lighter scales. We can compare the variance between the groups quite intuitively by colors.

```
# the mean value in total
total_avg = rfm.iloc[:, 0:3].mean()
total_avg
# calculate the proportional gap with total mean
cluster_avg = rfm.groupby('RFM_Level').mean().iloc[:, 0:3]
prop_rfm = cluster_avg/total_avg - 1
# heatmap with RFM
sns.heatmap(prop_rfm, cmap= 'Oranges', fmt= '.2f', annot = True)
plt.title('Heatmap of RFM quantile')
plt.plot()
```

And then repeat the same code for K-clusters as we did before.

```
# calculate the proportional gap with total mean
cluster_avg_K = rfm.groupby('K_Cluster').mean().iloc[:, 0:3]
prop_rfm_K = cluster_avg_K/total_avg - 1
# heatmap with K-means
sns.heatmap(prop_rfm_K, cmap= 'Blues', fmt= '.2f', annot = True)
plt.title('Heatmap of K-Means')
plt.plot()
```



It could be seen unmatching, especially at the top of the plots. But It's just because of the different order. The **Green** group on the left will correspond to group **2**. If you see the values inside each box, you can see the difference between the groups become significant for **Gold** and **1** group. And it could be easily recognized by the darkness of the color.

Programming  
Natural Language Processing  
Computer Vision  
News  
SQL


[JOIN NEWSLETTER](#)

This also tells us on which customer to focus on and to whom give special offers or promotions for fostering loyalty among customers. We can select the best communication channel for each segment and improve new marketing strategies.

### Resources

- A nice article about RFM analysis: <https://clevertap.com/blog/rfm-analysis/>
- Another useful explanation for RFM analysis: <https://www.optimove.com/learning-center/rfm-segmentation>
- Intuitive explanation on K-means clustering: [https://www.youtube.com/watch?v=\\_aWzGGRcic](https://www.youtube.com/watch?v=_aWzGGRcic)

**Bio:** Jiwon Jeong is a Graduate Research Assistant at Yonsei University.


Original. Reposted with permission.

### Related:

- Clustering Using K-means Algorithm
- K-Means & Other Clustering Algorithms: A Quick Intro with Python
- Beginner's Guide to Customer Segmentation

### What do you think?

5 Responses



0 Comments   KDnuggets   [Disqus' Privacy Policy](#)

[Login](#)

[Favorite](#)   [Tweet](#)   [Share](#)

Sort by Best



Start the discussion...

LOG IN WITH



OR SIGN UP WITH DISQUS (?)

Name

Be the first to comment.

[Subscribe](#)   [Add Disqus to your site](#)   [Do Not Sell My Data](#)

### More On This Topic

- Real Data, Big Impact: UChicago Students Work to Improve Sales at...
- KDnuggets™ News 19:n21, Jun 5: Transitioning your Career to...

SHARES

Programming  
Natural Language Processing  
Computer Vision  
News  
SQL

[JOIN NEWSLETTER](#)

Datasets  
Courses  
Online  
Certificates  
Webinars

Resources  
Meetings  
Jobs  
Software



Get KDnuggets, a leading newsletter on AI, Data Science and Machine Learning, straight to your inbox.

[SIGN UP](#)

By subscribing you accept KDnuggets Privacy Policy

Pages: 1 2 |

[<= Previous post](#)[Next post =>](#)

## Top Posts Past 30 Days

- 1 [Is Data Science a Dying Career?](#)
- 2 [The High Paying Side Hustles for Data Scientists](#)
- 3 [SQL Interview Questions for Experienced Professionals](#)
- 4 [Top Programming Languages and Their Uses](#)
- 5 [Top Five SQL Window Functions You Should Know For Data Science Interviews](#)
- 6 [A Deep Look Into 13 Data Scientist Roles and Their Responsibilities](#)
- 7 [7 Steps to Mastering Machine Learning with Python in 2022](#)
- 8 [Why Do Machine Learning Models Die In Silence?](#)
- 9 [Deliver a Killer Presentation in Data Science Interviews](#)
- 10 [Why are More Developers Using Python for Their Machine Learning Projects?](#)

© 2022 KDnuggets. | [About KDnuggets](#) | [Contact](#) | [Privacy policy](#) | [Terms of Service](#)

SHARES